

# How to Train Neural Networks

Seth D Temple, PhD  
Schmidt AI in Science Postdoc



# Agenda (exercises throughout)

1. Generate training data
  - a. Modularize simulation code
  - b. Data splitting
2. Architectures
  - a. Linear layers
  - b. Convolutional neural network
  - c. Recurrent neural networks
3. Neural networks as objects
4. Training loops
5. Evaluation

<https://github.com/sdtemple/zootopia3>

# Set up your environment

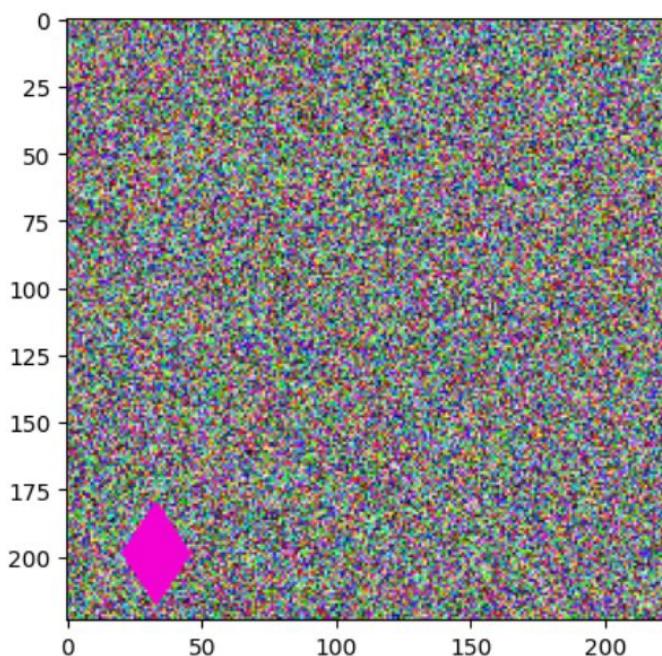
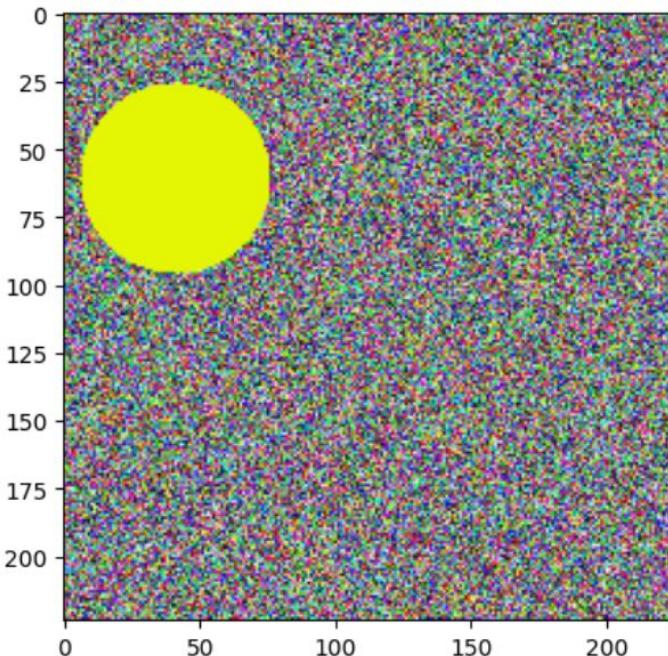
1. Fork the code at [github.com/sdtemple/zootopia3](https://github.com/sdtemple/zootopia3)
2. Move to IDE and terminal
  - a. git clone <https://github.com/youraccount/zootopia3.git>
  - b. Set up and activate environment  
(should be completed in tutorial 1)
3. Peruse the contents of the repository
  - a. `src/` : the Python package
  - b. `examples/` : where we will develop models
  - c. `data/` : where we will store code
4. Turn off AI coding assistant
  - a. In VS Code, Ctrl + , then search “editor.inlineSuggest.enabled”

# API references

- <https://docs.pytorch.org/docs/stable/pytorch-api.html>
- <https://numpy.org/doc/stable/reference/>
- [https://scikit-learn.org/stable/api/sklearn.model\\_selection.html](https://scikit-learn.org/stable/api/sklearn.model_selection.html)
- <https://scikit-learn.org/stable/api/sklearn.preprocessing.html>
- [https://matplotlib.org/stable/api/pyplot\\_summary.html](https://matplotlib.org/stable/api/pyplot_summary.html)

# Training data

The main task: classify shapes (4) and colors (6 - 8)



# Simulate training data

- [answers/simulate-exercise.ipynb](#)
- Modify the parameters to create data
  - Use different file suffixes for different data simulations
  - Start with a small number of simulated data examples
  - After getting a working model, scale up the data size

# Neural networks as objects

# Models via object-oriented programming

- Neural network is defined as a **Python class** that inherits from `torch.nn.Module`.
- **Encapsulation:** parameters, layers, and forward logic are bundled together
  - Modularity and reusability
- `__init__`: defines layers and components.
- `__forward__`: defines how data flows through them.
- **Composability:** layers are objects  
→ combine them to build complex architectures.

# The `__init__` function

- Declares what the object is!
- Constructor: runs when a model object is created
- Defines and stores the model's components and hyperparameters
- Establishes the internal state of the model

## `super()` within the method

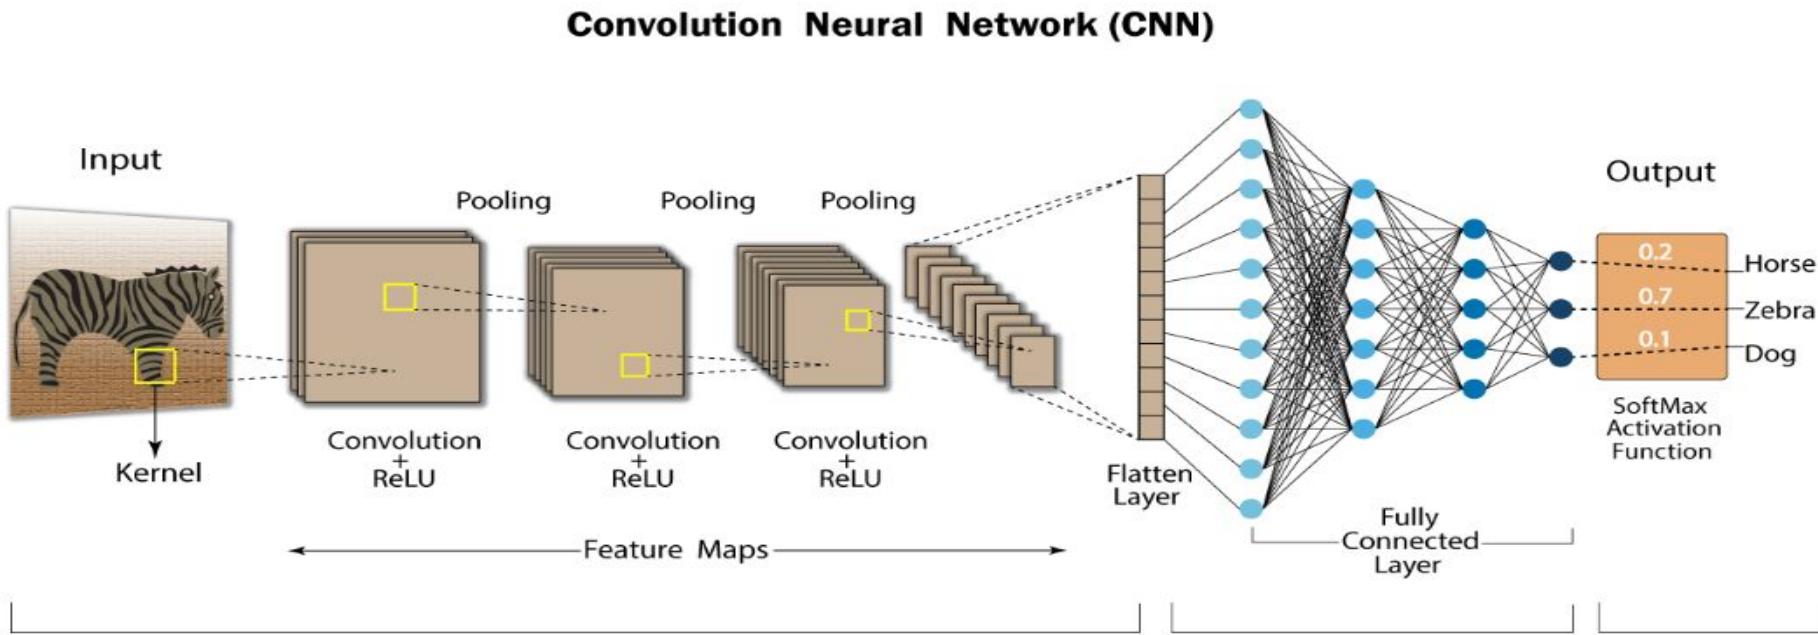
- Calls the parent class constructor `torch.nn.Module`
- Ensures inherited functionality is properly initialized
- Required for framework features (e.g., parameter tracking)

# self

- class is a blueprint, but self stores the specific instance
- **self.attribute =**
  - Stores data on the object
  - Makes it accessible to other methods
- You always pass self into the class methods!
  - def forward(self, x):
  - def do\_something(self, \*args, \*\*kwargs)
- Define attributes for the model architecture
  - **self.conv1 = nn.Conv2d()**
  - **self.linear1 = nn.Linear()**

# The forward(self, x) function

- Specifies the computation graph
  - How inputs are transformed into outputs
  - By invoking the self layer attributes
  - This is called when you run `model(x)`

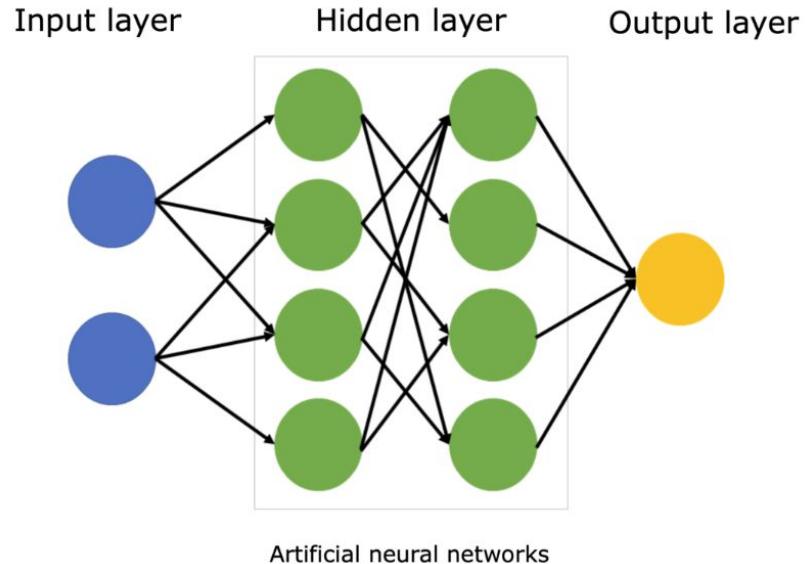


# Architectures

[examples/modeling-exercise-1.ipynb](#)  
[examples/modeling-exercise-2.ipynb](#)

# Linear() layer

- All-to-all mapping
  - The most amount of parameters!
  - Don't make this too large
- Dropout and/or batch normalization possible for all but last layer!
- Have to specify the first dimension, which changes in CNNs

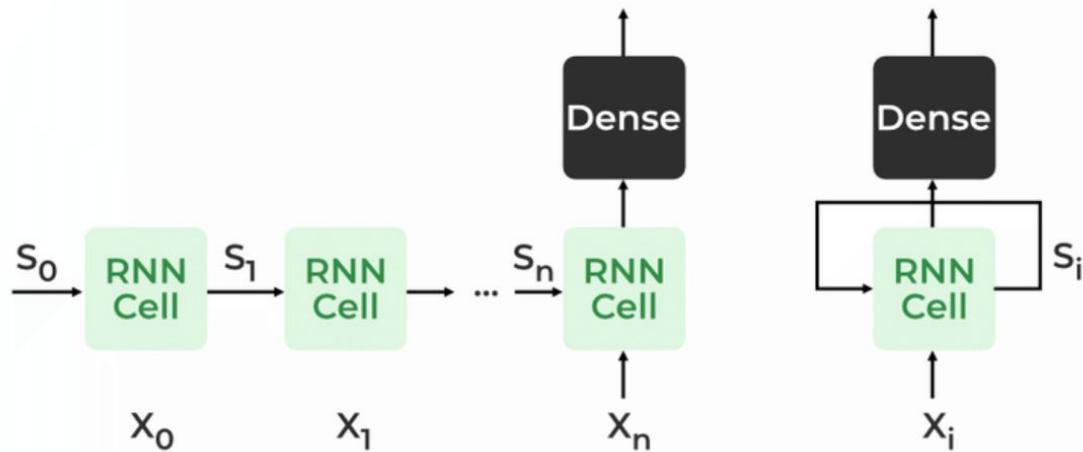


$$H_{out} = \left\lfloor \frac{H_{in} + 2 \times \text{padding}[0] - \text{dilation}[0] \times (\text{kernel\_size}[0] - 1) - 1}{\text{stride}[0]} + 1 \right\rfloor$$

# RNN() or LSTM() layer

- For sequential data
- Slower to train
- Careful about initializing hidden and cell states
- Careful about input dimensions

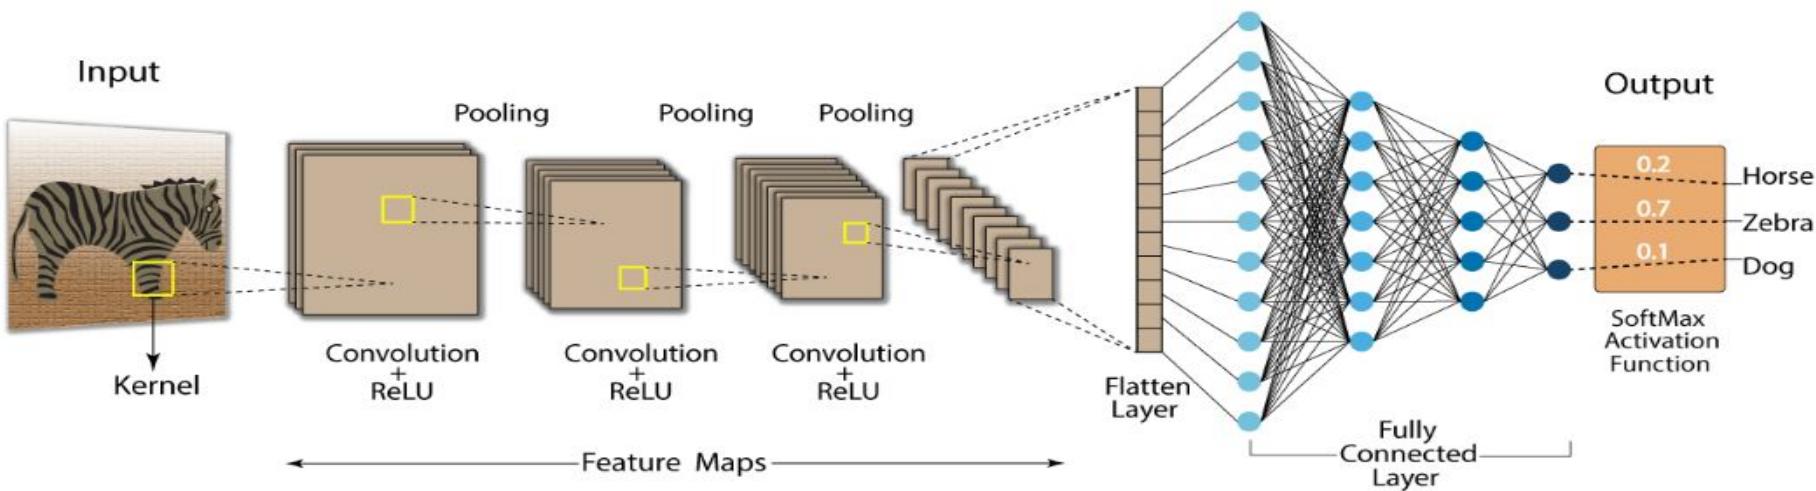
## RECURRENT NEURAL NETWORKS



# Conv2d() layer

- Input channels (e.g., RGB colors)
- Output channels
  - You often increase from initial (3) channels
  - For instance, 3 → 32 → 32 → Linear()
- **kernel\_size, stride, padding**
- Pooling and normalization layers follows

## Convolution Neural Network (CNN)



Source Layer

5	2	6	8	2	0	1	2
4	3	4	5	1	9	6	3
3	9	2	4	7	7	6	9
1	3	4	6	8	2	2	1
8	4	6	2	3	1	8	8
5	8	9	0	1	0	2	3
9	2	6	6	3	6	2	1
9	8	8	2	6	3	4	5

Convolutional Kernel

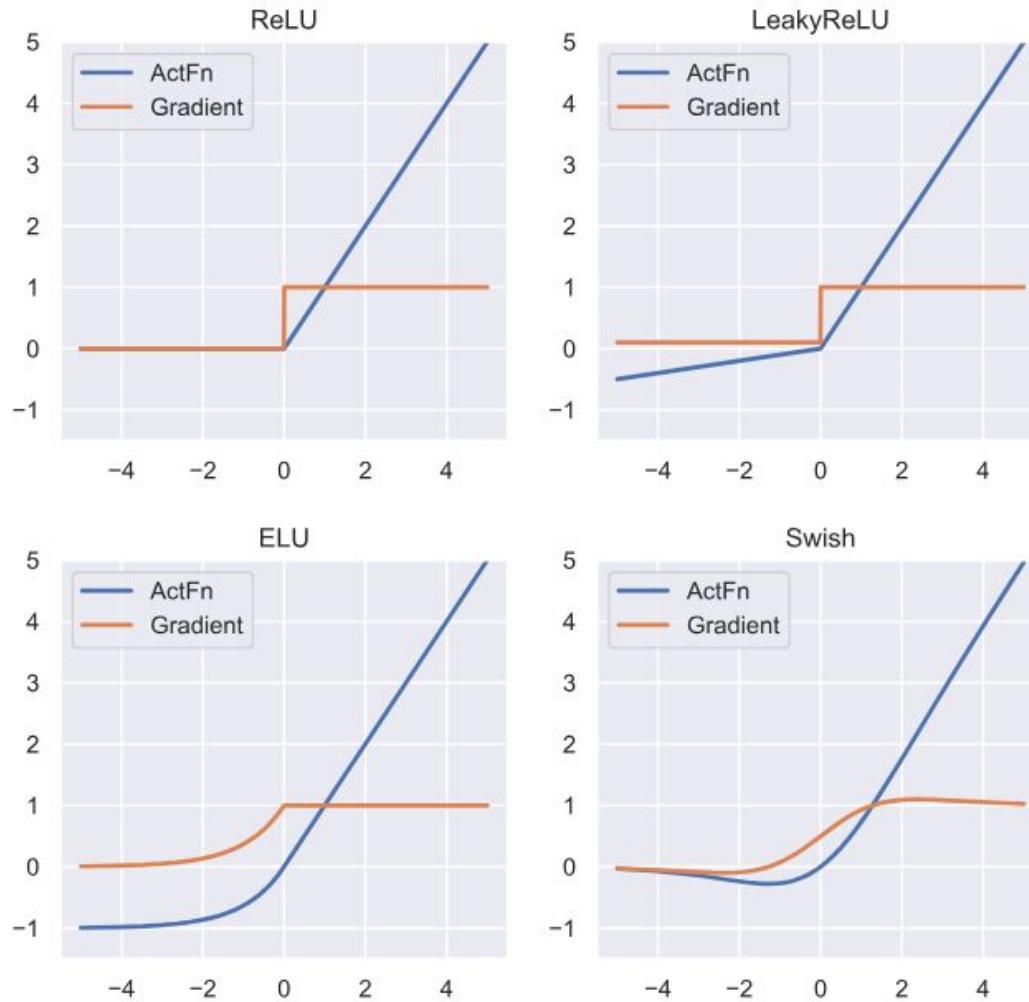
-1	0	1
2	1	2
1	-2	0

Destination Layer

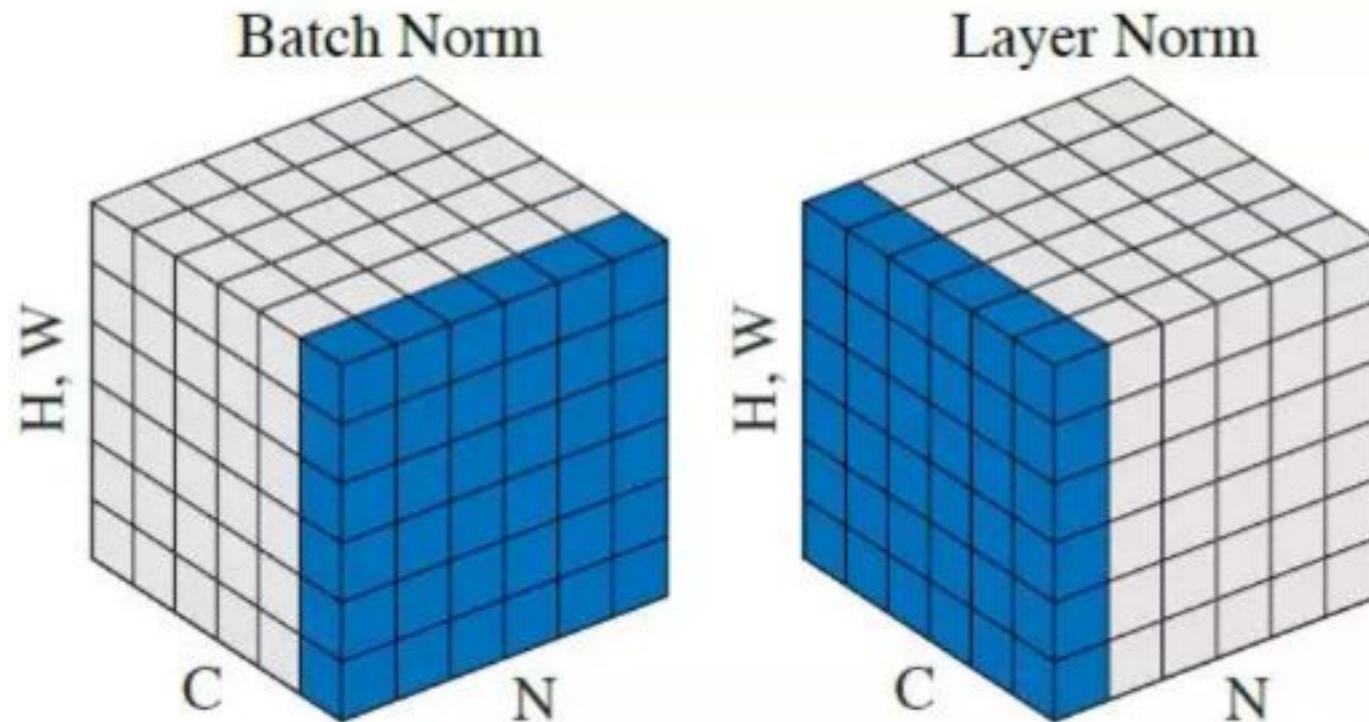

$$\begin{aligned} & (-1 \times 5) + (0 \times 2) + (1 \times 6) + \\ & (2 \times 4) + (1 \times 3) + (2 \times 4) + \\ & (1 \times 3) + (-2 \times 9) + (0 \times 2) = 5 \end{aligned}$$

# Other layers

- Pooling (same math as `Conv2d`):
  - `MaxPool2d()`
  - `AvgPool2d()`
- Activations
  - <https://docs.pytorch.org/docs/stable/nn.functional.html#non-linear-activation-functions>
- `Dropout()`
  - This is important to make inferences generalizable past training
- `LayerNorm1d()`
- `BatchNorm2d()`
  - Normalization layers can be super important!
  - Often your model will crash without them



# Normalization visualization



# Training loop

# Standard way to define your training loop

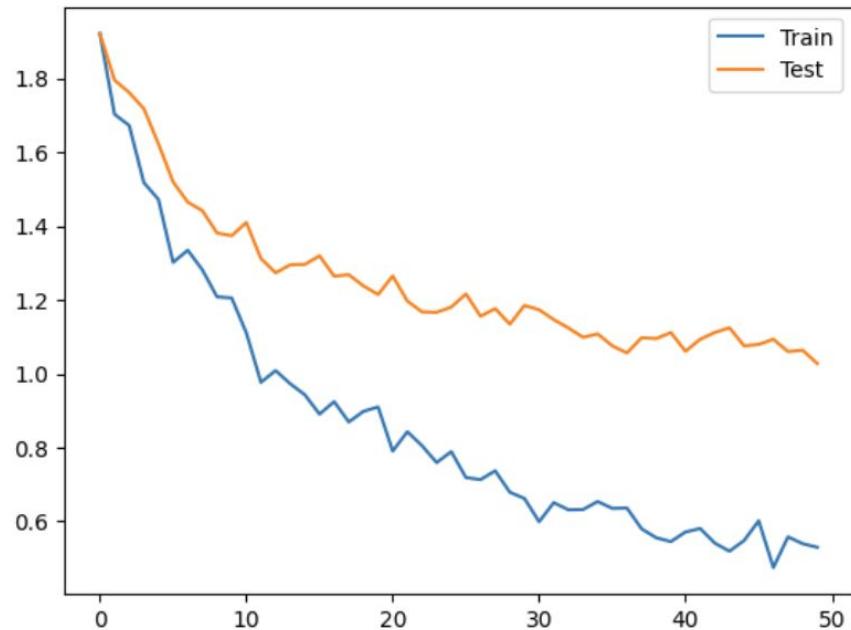
1. for epoch in range(num\_epochs)
  - a. How many times to pass over the data
2. **model.train()** : turns on the learning state!
3. for x, y in train\_loader:
4. **optimizer.zero\_grad()**
5. Invoke ` yhat = model(x) `
6. Invoke ` criterion(y, yhat)`
  - a. Care about dimensions of x and model output
  - b. Care about data being on the GPU device
7. **loss.backward()**
8. **optimizer.step()**

# Core training steps

- `zero_grad()` clears stored gradients from the previous iteration
- `loss.backward()` computes gradients via backpropagation
- `optimizer.step()` updates parameters using those gradients
- These steps implement one optimization iteration
- Order matters and is repeated every batch

# Diagnosing issues

- Loss not decreasing
  - Print the per-epoch loss with a running loss from batches
- Training–validation gap
  - Train - test - and validate
  - Keep track of validation loss in `model.eval()` state
- Exploding / vanishing gradients
  - This number should not be too small or large  
``print(next(model.parameters()).grad.abs().max())``



# Evaluation (on validation data)

# Metrics for categorical classifier

- Accuracy: overall fraction of correct predictions
- Confusion matrix: per-class error structure

[https://en.wikipedia.org/wiki/Confusion\\_matrix](https://en.wikipedia.org/wiki/Confusion_matrix)

- Precision & recall: false positives vs. negatives
- ROC–AUC / PR–AUC:  
threshold-independent performance

## Confusion Matrix

	Actually Positive (1)	Actually Negative (0)
Predicted Positive (1)	True Positives (TPs)	False Positives (FPs)
Predicted Negative (0)	False Negatives (FNs)	True Negatives (TNs)

# Code snippet, or there is one in the notebooks

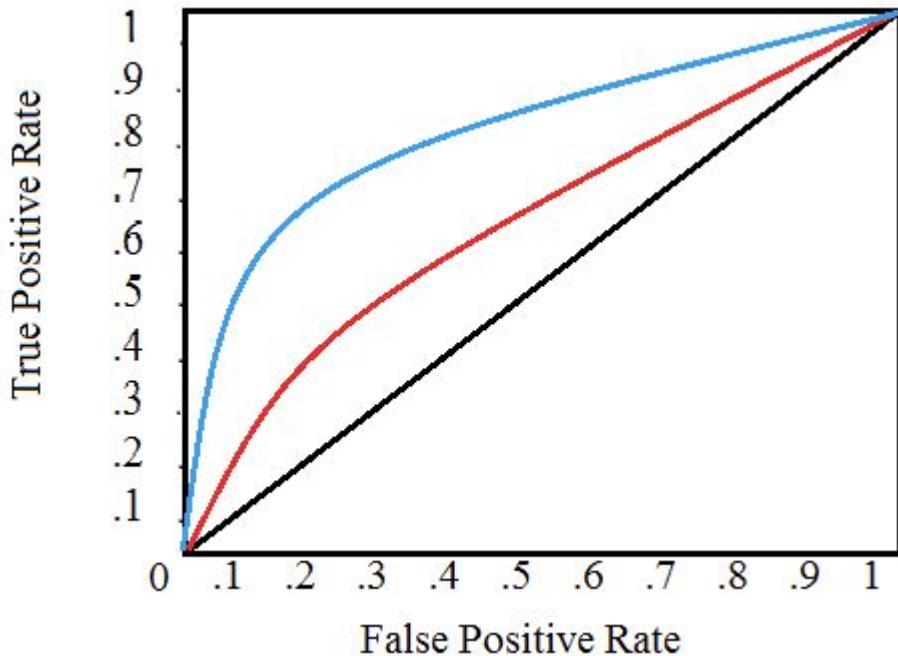
```
import matplotlib.pyplot as plt

def plot_confusion_matrix(df_confusion, title='Confusion matrix', cmap=plt.cm.gray_r):
    plt.matshow(df_confusion, cmap=cmap) # imshow
    #plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(df_confusion.columns))
    plt.xticks(tick_marks, df_confusion.columns, rotation=45)
    plt.yticks(tick_marks, df_confusion.index)
    #plt.tight_layout()
    plt.ylabel(df_confusion.index.name)
    plt.xlabel(df_confusion.columns.name)

df_confusion = pd.crosstab(y_actu, y_pred)
plot_confusion_matrix(df_confusion)
```

# Receiver operator curve

- TPR & FPR as we **change hard threshold** of a probabilistic classifier
- Prediction line headed to upper left is good
- Similar idea for PR curve (Precision-Recall)



Demo: train big model  
on the cluster

# The End

## Acknowledgement:

- Michigan Institute for Data & AI in Society
- Schmidt Sciences

