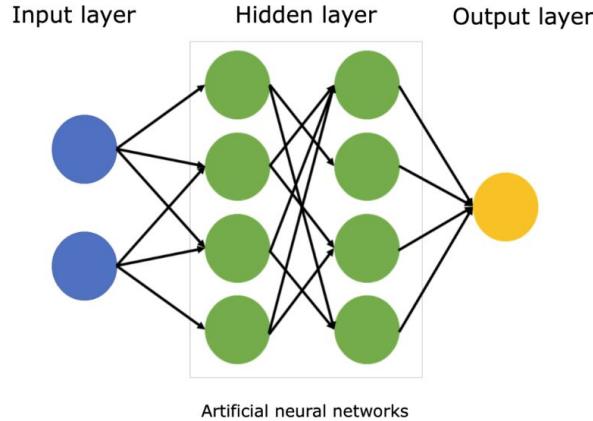


How to Train Neural Networks



Seth D Temple, PhD
Schmidt AI in Science Postdoc



Agenda (exercises throughout)

1. Generate training data
 - a. Modularize simulation code
 - b. Data splitting
2. Architectures
 - a. Linear layers
 - b. Convolutional neural network
 - c. Recurrent neural networks
3. Neural networks as objects
4. Training loops
5. Evaluation

<https://github.com/sdtemple/zootopia3>

Set up your environment

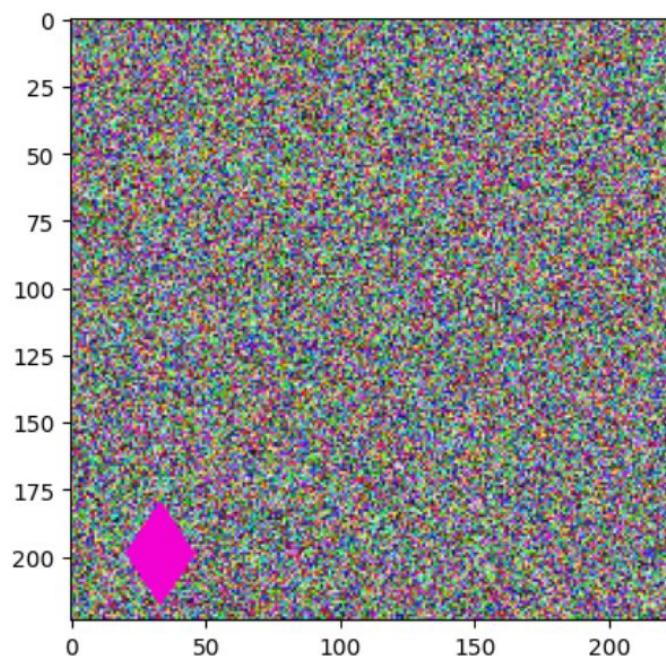
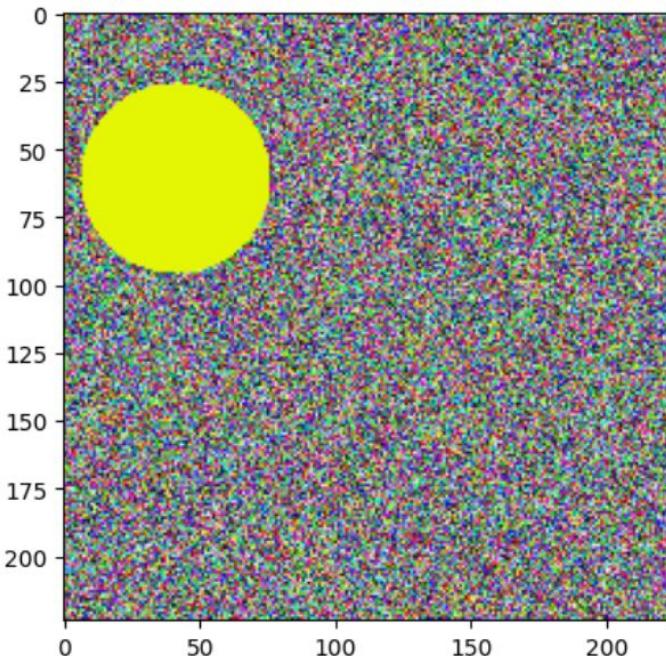
1. Fork the code at github.com/sdtemple/zootopia3
2. Move to IDE and terminal
 - a. git clone <https://github.com/youraccount/zootopia3.git>
 - b. Set up and activate environment
(should be completed in tutorial 1)
3. Peruse the contents of the repository
 - a. `src/` : the Python package
 - b. `examples/` : where we will develop models
 - c. `data/` : where we will store code
4. Turn off AI coding assistant
 - a. In VS Code, Ctrl + , then search “editor.inlineSuggest.enabled”

API references

- <https://docs.pytorch.org/docs/stable/pytorch-api.html>
- <https://numpy.org/doc/stable/reference/>
- https://scikit-learn.org/stable/api/sklearn.model_selection.html
- <https://scikit-learn.org/stable/api/sklearn.preprocessing.html>
- https://matplotlib.org/stable/api/pyplot_summary.html

Training data

The main task: classify shapes (4) and colors (6 - 8)



I wrote some redundant code in
examples/simulate-exercise.ipynb

- Separation of concerns: data, model, training, evaluation
- Easier debugging and faster iteration
- Reuse across experiments and projects
- Critical for scaling from notebooks to production

Modularize the examples/simulate.ipynb code

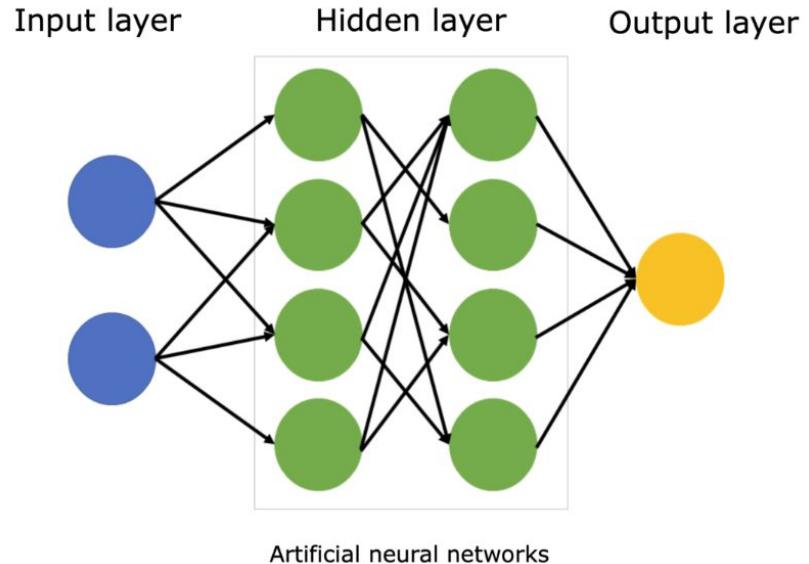
1. Skim the contents of [src/zootopia3/shapes](#)
2. Identify redundant features in for loop
3. Edit the `simulate_shapes()` function
in [src/zootopia3/simulate.py](#)
4. Now, edit the [examples/simulate.ipynb](#) file
5. Play around with the parameters in the notebook

Architectures

[examples/modeling-exercise-1.ipynb](#)
[examples/modeling-exercise-2.ipynb](#)

Linear() layer

- All-to-all mapping
 - The most amount of parameters!
 - Don't make this too large
- Dropout and/or batch normalization possible for all but last layer!
- Have to specify the first dimension, which changes in CNNs

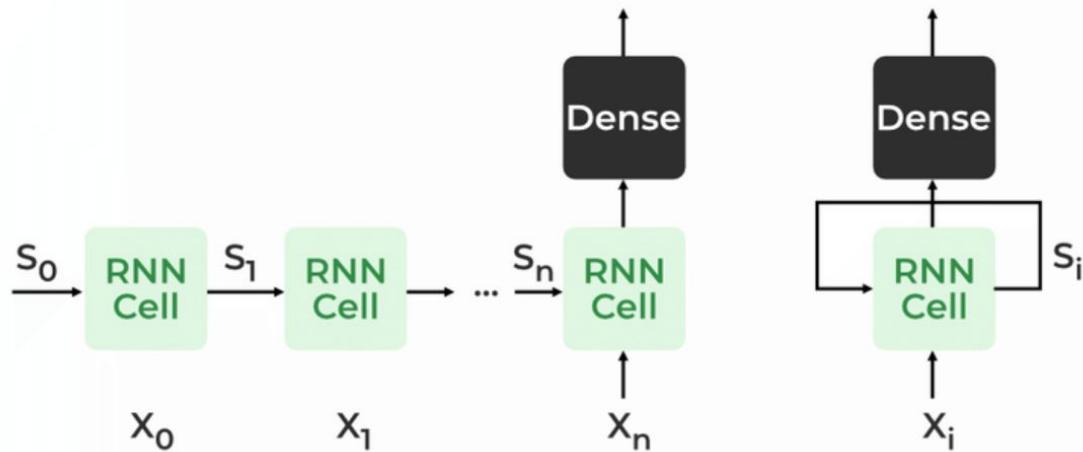


$$H_{out} = \left\lfloor \frac{H_{in} + 2 \times \text{padding}[0] - \text{dilation}[0] \times (\text{kernel_size}[0] - 1) - 1}{\text{stride}[0]} + 1 \right\rfloor$$

RNN() or LSTM() layer

- For sequential data
- Slower to train
- Careful about initializing hidden and cell states
- Careful about input dimensions

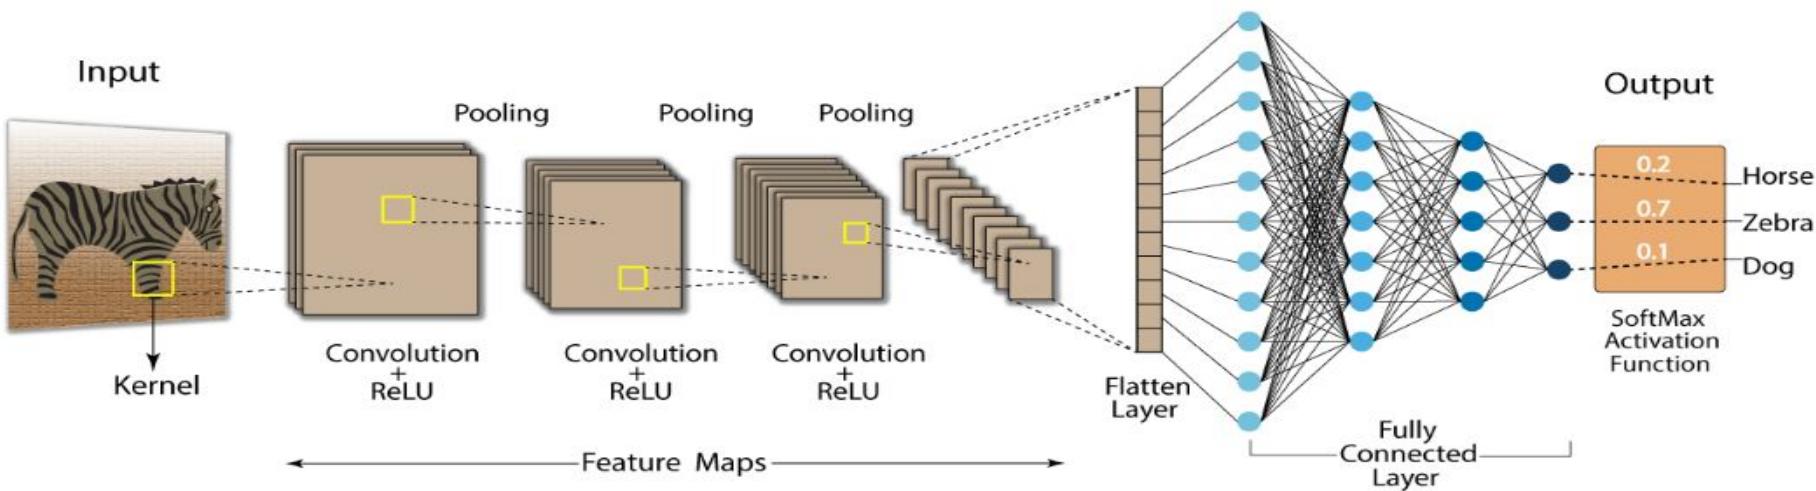
RECURRENT NEURAL NETWORKS



Conv2d() layer

- Input channels (e.g., RGB colors)
- Output channels
 - You often increase from initial (3) channels
 - For instance, 3 → 32 → 32 → Linear()
- **kernel_size, stride, padding**
- Pooling and normalization layers follows

Convolution Neural Network (CNN)



Source Layer

5	2	6	8	2	0	1	2
4	3	4	5	1	9	6	3
3	9	2	4	7	7	6	9
1	3	4	6	8	2	2	1
8	4	6	2	3	1	8	8
5	8	9	0	1	0	2	3
9	2	6	6	3	6	2	1
9	8	8	2	6	3	4	5

Convolutional Kernel

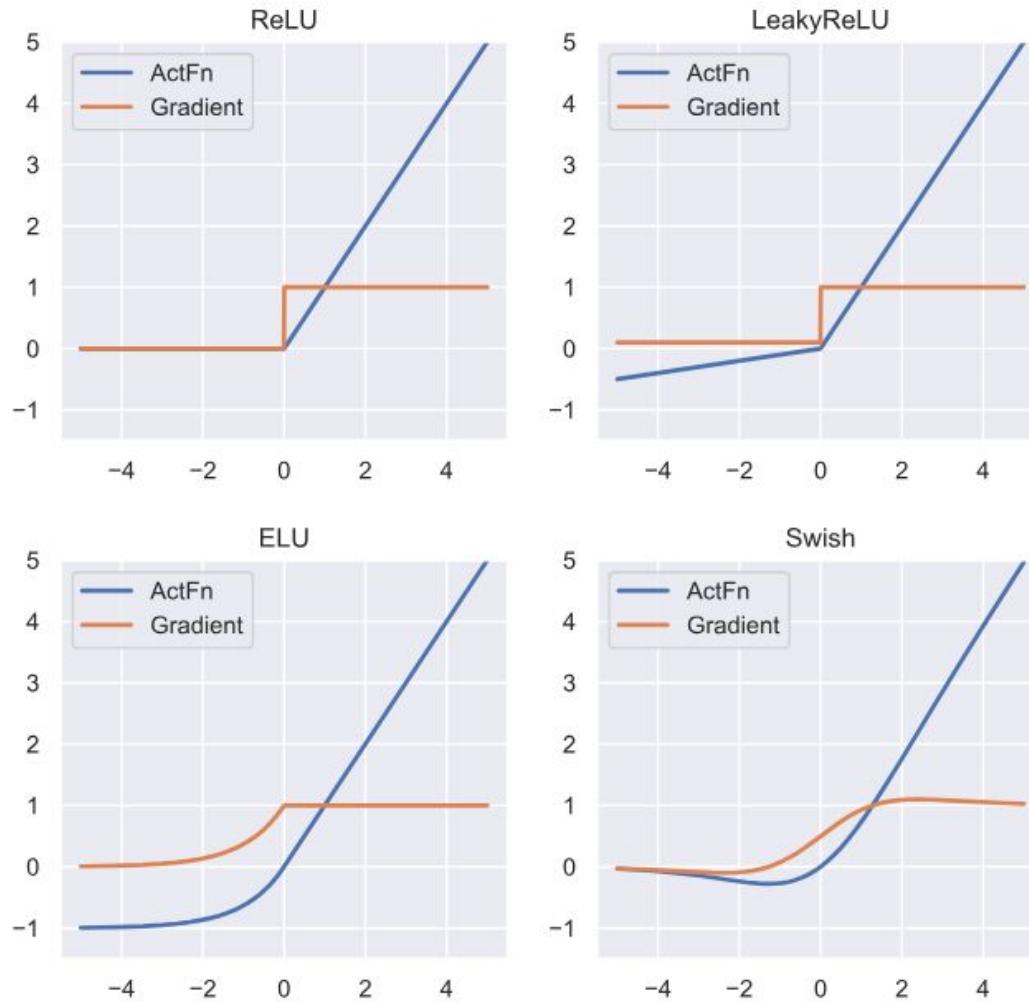
-1	0	1
2	1	2
1	-2	0

Destination Layer

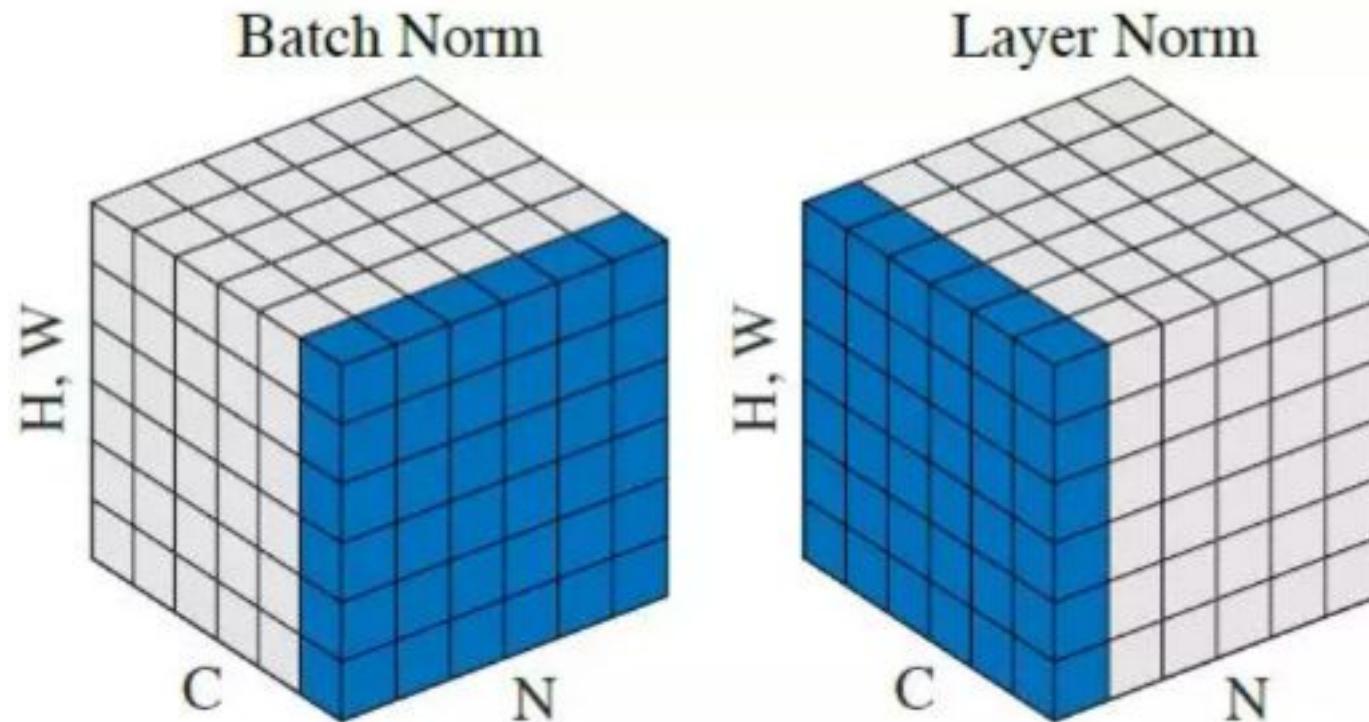
$$\begin{aligned} & (-1 \times 5) + (0 \times 2) + (1 \times 6) + \\ & (2 \times 4) + (1 \times 3) + (2 \times 4) + \\ & (1 \times 3) + (-2 \times 9) + (0 \times 2) = 5 \end{aligned}$$

Other layers

- Pooling (same math as `Conv2d`):
 - `MaxPool2d()`
 - `AvgPool2d()`
- Activations
 - <https://docs.pytorch.org/docs/stable/nn.functional.html#non-linear-activation-functions>
- `Dropout()`
 - This is important to make inferences generalizable past training
- `LayerNorm1d()`
- `BatchNorm2d()`
 - Normalization layers can be super important!
 - Often your model will crash without them



Normalization visualization



Neural networks as objects

Models via object-oriented programming

- Neural network is defined as a **Python class** that inherits from `torch.nn.Module`.
- **Encapsulation:** parameters, layers, and forward logic are bundled together
 - Modularity and reusability
- `__init__`: defines layers and components.
- `__forward__`: defines how data flows through them.
- **Composability:** layers are objects
→ combine them to build complex architectures.

The `__init__` function

- Declares what the object is!
- Constructor: runs when a model object is created
- Defines and stores the model's components and hyperparameters
- Establishes the internal state of the model

`super()` within the method

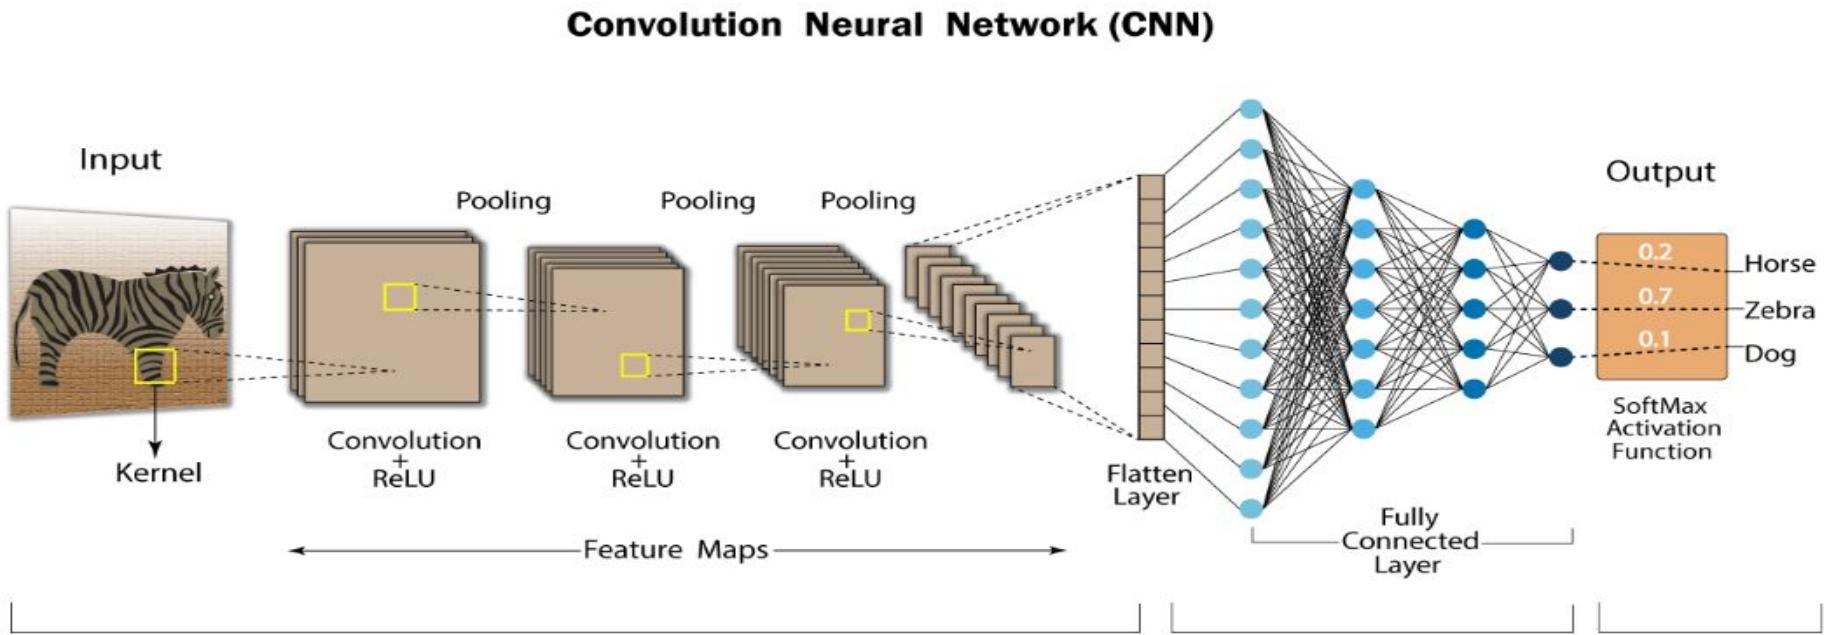
- Calls the parent class constructor `torch.nn.Module`
- Ensures inherited functionality is properly initialized
- Required for framework features (e.g., parameter tracking)

self

- class is a blueprint, but self stores the specific instance
- **self.attribute =**
 - Stores data on the object
 - Makes it accessible to other methods
- You always pass self into the class methods!
 - def forward(self, x):
 - def do_something(self, *args, **kwargs)
- Define attributes for the model architecture
 - **self.conv1 = nn.Conv2d()**
 - **self.linear1 = nn.Linear()**

The forward(self, x) function

- Specifies the computation graph
 - How inputs are transformed into outputs
 - By invoking the self layer attributes
 - This is called when you run `model(x)`



Training loop

Standard way to define your training loop

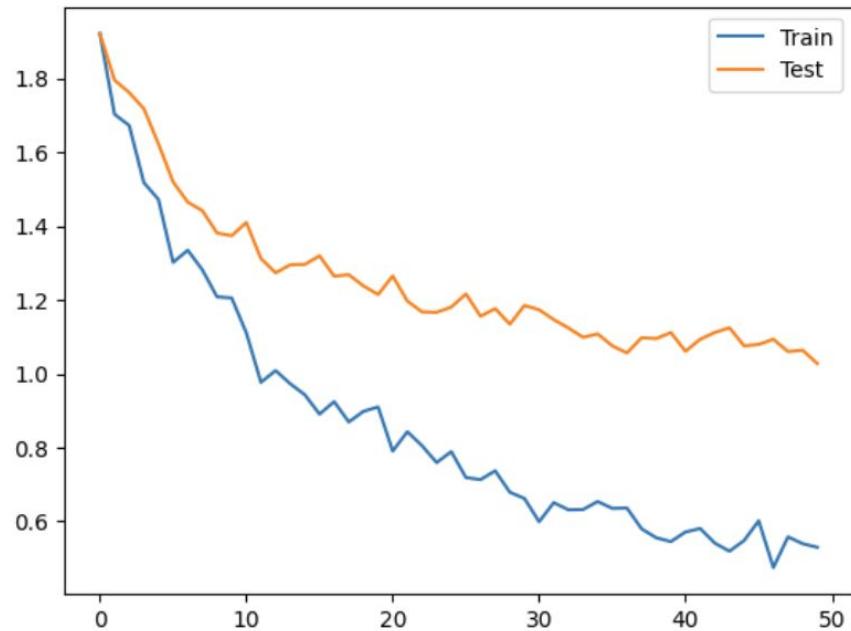
1. for epoch in range(num_epochs)
 - a. How many times to pass over the data
2. **model.train()** : turns on the learning state!
3. for x, y in train_loader:
4. **optimizer.zero_grad()**
5. Invoke ` yhat = model(x) `
6. Invoke ` criterion(y, yhat)`
 - a. Care about dimensions of x and model output
 - b. Care about data being on the GPU device
7. **loss.backward()**
8. **optimizer.step()**

Core training steps

- `zero_grad()` clears stored gradients from the previous iteration
- `loss.backward()` computes gradients via backpropagation
- `optimizer.step()` updates parameters using those gradients
- These steps implement one optimization iteration
- Order matters and is repeated every batch

Diagnosing issues

- Loss not decreasing
 - Print the per-epoch loss with a running loss from batches
- Training–validation gap
 - Train - test - and validate
 - Keep track of validation loss in `model.eval()` state
- Exploding / vanishing gradients
 - This number should not be too small or large
``print(next(model.parameters()).grad.abs().max())``



Evaluation (on validation data)

Metrics for categorical classifier

- Accuracy: overall fraction of correct predictions
- Confusion matrix: per-class error structure

https://en.wikipedia.org/wiki/Confusion_matrix

- Precision & recall: false positives vs. negatives
- ROC–AUC / PR–AUC:
threshold-independent performance

Confusion Matrix

	Actually Positive (1)	Actually Negative (0)
Predicted Positive (1)	True Positives (TPs)	False Positives (FPs)
Predicted Negative (0)	False Negatives (FNs)	True Negatives (TNs)

Code snippet, or there is one in the notebooks

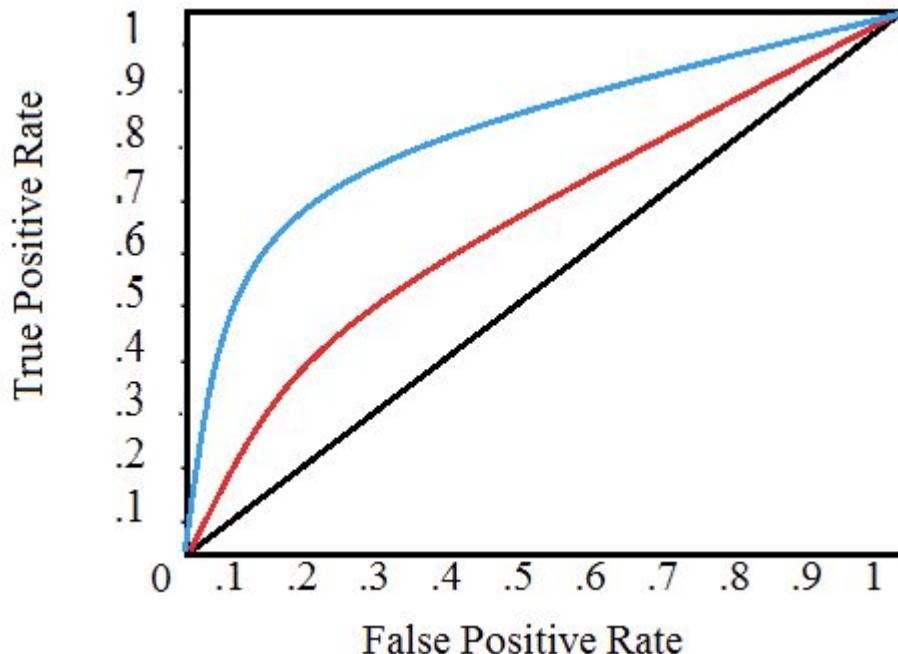
```
import matplotlib.pyplot as plt

def plot_confusion_matrix(df_confusion, title='Confusion matrix', cmap=plt.cm.gray_r):
    plt.matshow(df_confusion, cmap=cmap) # imshow
    #plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(df_confusion.columns))
    plt.xticks(tick_marks, df_confusion.columns, rotation=45)
    plt.yticks(tick_marks, df_confusion.index)
    #plt.tight_layout()
    plt.ylabel(df_confusion.index.name)
    plt.xlabel(df_confusion.columns.name)

df_confusion = pd.crosstab(y_actu, y_pred)
plot_confusion_matrix(df_confusion)
```

Receiver operator curve

- TPR & FPR as we **change hard threshold** of a probabilistic classifier
- Prediction line headed to upper left is good
- Similar idea for PR curve (Precision-Recall)



Conformal prediction

- **Prediction intervals** with finite-sample guarantees
 - The AI/ML crew's version of a confidence interval
- On top of any model (regression or classification)
- Nonparametric and distribution-free in principle
- Uses a **calibration set** to adjust predictions

(???) [examples/conformal-prediction.ipynb](#)

Demo: train big model
on the cluster

The End

Acknowledgement:

- Michigan Institute for Data & AI in Society
- Schmidt Sciences

